



Machine Learning

Abdelhak Mahmoudi
abdelhak.mahmoudi@um5.ac.ma

2020

Content

1. The Big Picture

2. Supervised Learning

- Linear Regression, Logistic Regression, Support Vector Machines, Trees, Random Forests, Boosting, Artificial Neural Networks

3. **Unsupervised Learning**


- Principal Component Analysis, K-means, Mean Shift

Unsupervised Learning

- **Dimensionality Reduction**
 - **Principal Component Analysis (PCA)**
- Clustering
 - K-Means
 - Mean-Shift



Dimensionality Reduction

- Curse of dimensionality ($n \gg m$)
 - Data are at risk of being **very sparse** in high dimensional space
 - High risk of **overfitting**




Dimensionality Reduction

- Transforms feature space from n to k ($k < n$)
 - Some **features** are probably **corelated** (dependent)
 - Some **features** are almost **constant**
 - **Transform but preserve** the maximum of **variance**



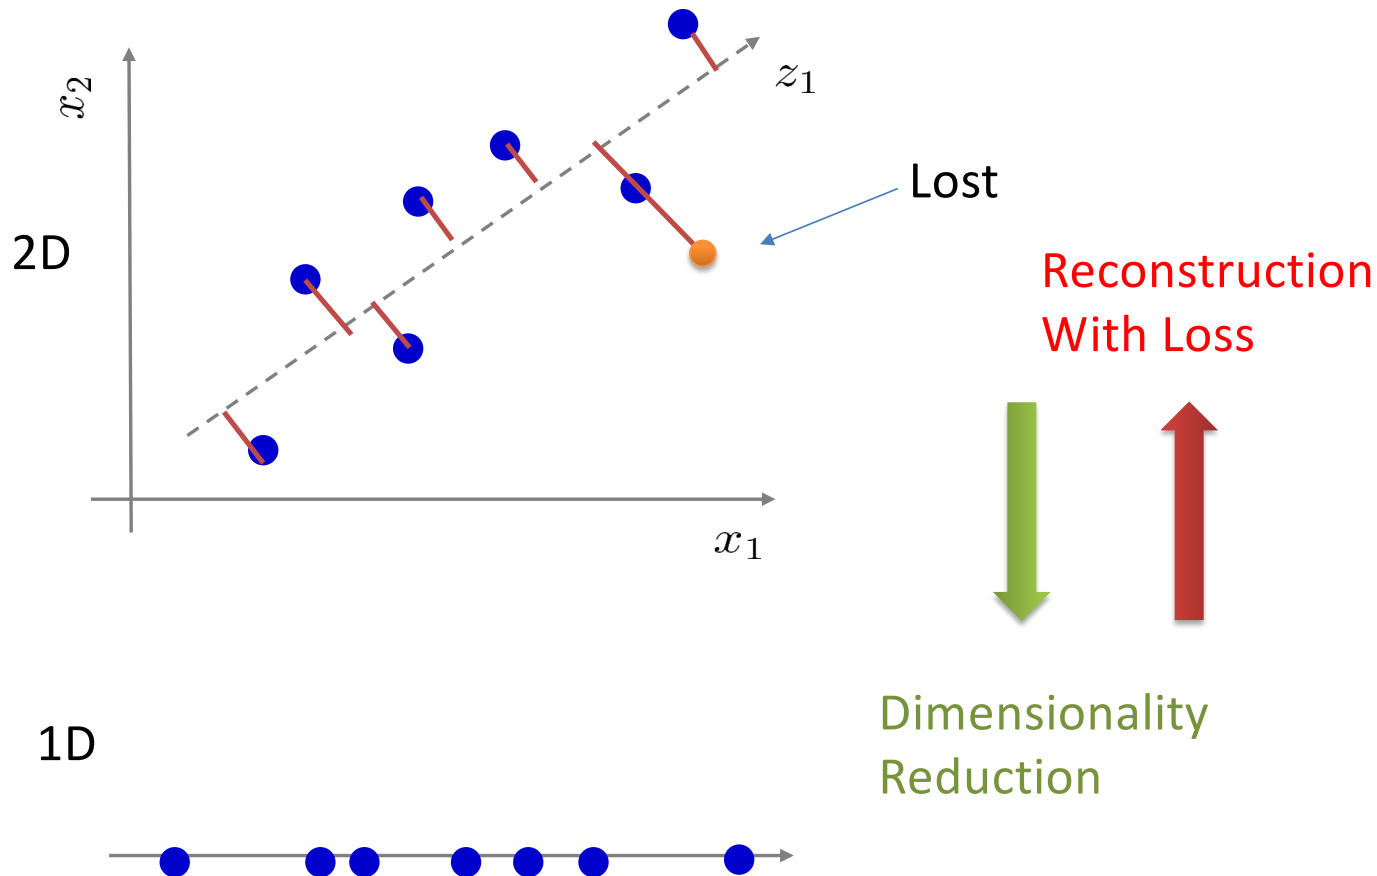


Dimensionality Reduction


- Often
 - **Not** necessarily lead to **better performance**
 - **Not the better** way to address **overfitting !**
- Always
 - **Speed up** training
 - Allow **data compression**
 - Allow **data exploration**
 - Allow **data visualization** (DataViz)




Principal Component Analysis



Principal Component Analysis



Maximum loss
Less variance



Minimum loss
More variance

Principal Component Analysis

- Singular Value Decomposition (SVD) (very costly)
 - Parallelization: Incremental PCA (fast), Randomized PCA (faster)
- PCA assumes that the dataset is centered around the origin
- How many dimensions to preserve?
 - Reduce dimensions that add up to a sufficiently large portion of the explained variance (e.g., 99%)
- Kernel PCA (kPCA): use the kernel trick like SVM
- In practice, use kPCA to transform the feature space, then perform classification or regression or clustering.

Principal Component Analysis

- Hyper-Parameters Tuning
 - d : polynomial Kernel
 - γ : RBF kernel
 - k : Number of retained principal components
 - Etc.

Unsupervised Learning


- Dimensionality Reduction
 - Principal Component Analysis (PCA)
- **Clustering**
 - **K-Means**
 - **Mean-Shift**

K-Means

- Pick a number of clusters **k**
- Initialize **centroids** randomly
- Problem of **local optima**
 - Run K-means a lot of times
- **Sensible** to **initial** conditions
- Have to **specify k** !

Repeat until convergence:

Assign each example to the cluster of the **nearest** centroid
Compute the **mean** in each cluster
Put the mean as the **new centroid**




Mean Shift

- Chose a **radius r of the clusters**
- Initialize **centroids** at each example
- **No need** to specify the number of clusters

For each example:

Repeat until convergence:

Compute the **mean** in its cluster with radius r
Shift the cluster to the new mean centroid



Other Clustering methods

- Expectation Maximization (EM)
- Hierarchical Clustering
- Affinity Propagation (AP)
- Etc.